

# SNS\_ASS5

April 8, 2021

```
[1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from numpy import e, log
import seaborn as sns
from itertools import combinations
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import OrdinalEncoder
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score

from sklearn import svm

# import warnings
import warnings
# filter warnings
warnings.filterwarnings('ignore')
```

```
[2]: from google.colab import drive
drive.mount("/content/drive/")
```

Mounted at /content/drive/

```
[3]: df = pd.read_csv("drive/My Drive/full.csv")

print("Number of data points:", df.shape[0])
```

Number of data points: 4898430

```
[4]: df.head()
```

```
[4]:  0  tcp  http  SF  215  ...  0.00.10  0.00.11  0.00.12  0.00.13  normal.
      0  0  tcp  http  SF  162  ...      0.0      0.0      0.0      0.0  normal.
      1  0  tcp  http  SF  236  ...      0.0      0.0      0.0      0.0  normal.
      2  0  tcp  http  SF  233  ...      0.0      0.0      0.0      0.0  normal.
      3  0  tcp  http  SF  239  ...      0.0      0.0      0.0      0.0  normal.
      4  0  tcp  http  SF  238  ...      0.0      0.0      0.0      0.0  normal.
```

[5 rows x 42 columns]

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4898430 entries, 0 to 4898429
Data columns (total 42 columns):
 #   Column  Dtype
---  -
 0    0      int64
 1   tcp      object
 2   http     object
 3    SF      object
 4   215      int64
 5  45076     int64
 6   0.1      int64
 7   0.2      int64
 8   0.3      int64
 9   0.4      int64
10  0.5      int64
11  1        int64
12  0.6      int64
13  0.7      int64
14  0.8      int64
15  0.9      int64
16  0.10     int64
17  0.11     int64
18  0.12     int64
19  0.13     int64
20  0.14     int64
21  0.15     int64
22  1.1      int64
23  1.2      int64
24  0.00     float64
25  0.00.1   float64
26  0.00.2   float64
27  0.00.3   float64
28  1.00     float64
29  0.00.4   float64
30  0.00.5   float64
31  0.16     int64
```

```

32  0.17      int64
33  0.00.6    float64
34  0.00.7    float64
35  0.00.8    float64
36  0.00.9    float64
37  0.00.10   float64
38  0.00.11   float64
39  0.00.12   float64
40  0.00.13   float64
41  normal.   object
dtypes: float64(15), int64(23), object(4)
memory usage: 1.5+ GB

```

### 0.0.1 Object data types:

It is clear from the above observation that there are four columns whose data types are object type i.e, categorical features and we need to encode them to numerical type. So, we will use some library to encode them to numerical.

```
[6]: obj_df = df.select_dtypes(include=['object']).copy()
obj_df.head()
```

```
[6]:   tcp  http  SF  normal.
0  tcp  http  SF  normal.
1  tcp  http  SF  normal.
2  tcp  http  SF  normal.
3  tcp  http  SF  normal.
4  tcp  http  SF  normal.
```

```
[7]: obj_df[obj_df.isnull().any(axis=1)]
```

```
[7]: Empty DataFrame
Columns: [tcp, http, SF, normal.]
Index: []
```

```
[8]: obj_df["tcp"].value_counts()
```

```
[8]: icmp      2833545
tcp        1870597
udp         194288
Name: tcp, dtype: int64
```

```
[9]: obj_df["http"].value_counts()
```

```
[9]: ecr_i      2811660
private     1100831
http        623090
smtp        96554
other       72653
...
tftp_u      3
aol         2
```

```
harvest          2
http_8001        2
http_2784        1
Name: http, Length: 70, dtype: int64
```

```
[10]: obj_df["SF"].value_counts()
```

```
[10]: SF          3744327
      S0          869829
      REJ         268874
      RSTR          8094
      RST0         5344
      SH           1040
      S1           532
      S2           161
      RST0S0        122
      OTH           57
      S3           50
      Name: SF, dtype: int64
```

```
[11]: obj_df["normal."].value_counts()
```

```
[11]: smurf.          2807886
      neptune.       1072017
      normal.        972780
      satan.         15892
      ipsweep.       12481
      portsweep.     10413
      nmap.          2316
      back.          2203
      warezclient.   1020
      teardrop.      979
      pod.           264
      guess_passwd.  53
      buffer_overflow. 30
      land.          21
      warezmaster.   20
      imap.          12
      rootkit.       10
      loadmodule.    9
      ftp_write.     8
      multihop.      7
      phf.           4
      perl.          3
      spy.           2
      Name: normal., dtype: int64
```

```
[12]: ord_enc = OrdinalEncoder()
      df["make_tcp"] = ord_enc.fit_transform(df[["tcp"]])
```

```
df[["tcp", "make_tcp"]].head(11)
```

```
[12]:   tcp  make_tcp
0   tcp        1.0
1   tcp        1.0
2   tcp        1.0
3   tcp        1.0
4   tcp        1.0
5   tcp        1.0
6   tcp        1.0
7   tcp        1.0
8   tcp        1.0
9   tcp        1.0
10  tcp        1.0
```

```
[13]: ord_enc = OrdinalEncoder()
df["make_http"] = ord_enc.fit_transform(df[["http"]])
df[["http", "make_http"]].head(11)
```

```
[13]:   http  make_http
0   http        24.0
1   http        24.0
2   http        24.0
3   http        24.0
4   http        24.0
5   http        24.0
6   http        24.0
7   http        24.0
8   http        24.0
9   http        24.0
10  http        24.0
```

```
[14]: ord_enc = OrdinalEncoder()
df["make_SF"] = ord_enc.fit_transform(df[["SF"]])
df[["SF", "make_SF"]].head(11)
```

```
[14]:   SF  make_SF
0   SF        9.0
1   SF        9.0
2   SF        9.0
3   SF        9.0
4   SF        9.0
5   SF        9.0
6   SF        9.0
7   SF        9.0
8   SF        9.0
9   SF        9.0
10  SF        9.0
```

```
[15]: #ord_enc = OrdinalEncoder()
#df["make_normal"] = ord_enc.fit_transform(df[["normal."]])

#df[["normal.", "make_normal"]].head(11)
```

```
[16]: df.head(10)
```

```
[16]:    0  tcp  http  SF  215  ...  0.00.13  normal.  make_tcp  make_http  make_SF
0  0  tcp  http  SF  162  ...    0.0  normal.    1.0    24.0    9.0
1  0  tcp  http  SF  236  ...    0.0  normal.    1.0    24.0    9.0
2  0  tcp  http  SF  233  ...    0.0  normal.    1.0    24.0    9.0
3  0  tcp  http  SF  239  ...    0.0  normal.    1.0    24.0    9.0
4  0  tcp  http  SF  238  ...    0.0  normal.    1.0    24.0    9.0
5  0  tcp  http  SF  235  ...    0.0  normal.    1.0    24.0    9.0
6  0  tcp  http  SF  234  ...    0.0  normal.    1.0    24.0    9.0
7  0  tcp  http  SF  239  ...    0.0  normal.    1.0    24.0    9.0
8  0  tcp  http  SF  181  ...    0.0  normal.    1.0    24.0    9.0
9  0  tcp  http  SF  184  ...    0.0  normal.    1.0    24.0    9.0
```

[10 rows x 45 columns]

```
[17]: target=df['normal.'].values
df.drop(['tcp', 'http', 'SF', 'normal.'], axis=1, inplace=True)

df.head(10)
```

```
[17]:    0  215  45076  0.1  0.2  ...  0.00.12  0.00.13  make_tcp  make_http  make_SF
0  0  162   4528   0   0  ...    0.0    0.0    1.0    24.0    9.0
1  0  236   1228   0   0  ...    0.0    0.0    1.0    24.0    9.0
2  0  233   2032   0   0  ...    0.0    0.0    1.0    24.0    9.0
3  0  239    486   0   0  ...    0.0    0.0    1.0    24.0    9.0
4  0  238   1282   0   0  ...    0.0    0.0    1.0    24.0    9.0
5  0  235   1337   0   0  ...    0.0    0.0    1.0    24.0    9.0
6  0  234   1364   0   0  ...    0.0    0.0    1.0    24.0    9.0
7  0  239   1295   0   0  ...    0.0    0.0    1.0    24.0    9.0
8  0  181   5450   0   0  ...    0.0    0.0    1.0    24.0    9.0
9  0  184    124   0   0  ...    0.0    0.0    1.0    24.0    9.0
```

[10 rows x 41 columns]

```
[18]: target
```

```
[18]: array(['normal.', 'normal.', 'normal.', ..., 'normal.', 'normal.',
       'normal.'], dtype=object)
```

```
[19]: Y=target

X=df.values

print(X[0])
```

```
[0.000e+00 1.620e+02 4.528e+03 0.000e+00 0.000e+00 0.000e+00 0.000e+00
0.000e+00 1.000e+00 0.000e+00 0.000e+00 0.000e+00 0.000e+00 0.000e+00
0.000e+00 0.000e+00 0.000e+00 0.000e+00 0.000e+00 2.000e+00 2.000e+00
0.000e+00 0.000e+00 0.000e+00 0.000e+00 1.000e+00 0.000e+00 0.000e+00
1.000e+00 1.000e+00 1.000e+00 0.000e+00 1.000e+00 0.000e+00 0.000e+00
0.000e+00 0.000e+00 0.000e+00 1.000e+00 2.400e+01 9.000e+00]
```

```
[20]: Y[:10]
```

```
[20]: array(['normal.', 'normal.', 'normal.', 'normal.', 'normal.', 'normal.',
'normal.', 'normal.', 'normal.', 'normal.'], dtype=object)
```

### Limiting Data points:

```
[ ]: #X=X[:1500000]
#Y=Y[:1500000]
```

```
[ ]: from sklearn.decomposition import PCA
```

```
[ ]: #pca = PCA(n_components=10)
#X=pca.fit_transform(X)
```

```
[21]: X_train, X_val, y_train, y_val = train_test_split(X, Y, test_size=0.33,
↳random_state=42)
```

```
[22]: print(X_train.shape)
print(y_train.shape)

print(X_val.shape)
print(y_val.shape)
```

```
(3281948, 41)
(3281948,)
(1616482, 41)
(1616482,)
```

## 0.1 Testing related data

```
[23]: df2 = pd.read_csv("drive/My Drive/test.csv",header=None)

print("Number of data points:",df2.shape[0])
```

Number of data points: 311029

```
[24]: df2.head(10)
```

```
[24]:    0    1    2      3    4    ...    37    38    39    40    41
0  NaN    0  udp  private  SF  ...  0.00.7  0.00.8  0.00.9  0.00.10  0.00.11
1  0.0    0  udp  private  SF  ...    0.0    0.0    0.0    0.0    0.0
2  1.0    0  udp  private  SF  ...    0.0    0.0    0.0    0.0    0.0
```

3	2.0	0	udp	private	SF	...	0.0	0.0	0.0	0.0	0.0
4	3.0	0	udp	private	SF	...	0.0	0.0	0.0	0.0	0.0
5	4.0	0	udp	private	SF	...	0.0	0.0	0.0	0.0	0.0
6	5.0	0	udp	domain_u	SF	...	0.0	0.0	0.0	0.0	0.0
7	6.0	0	udp	private	SF	...	0.0	0.0	0.0	0.0	0.0
8	7.0	0	udp	private	SF	...	0.0	0.0	0.0	0.0	0.0
9	8.0	0	tcp	http	SF	...	0.01	0.0	0.0	0.0	0.0

[10 rows x 42 columns]

```
[25]: ### Dropping the first column as it is an index
```

```
df2.drop(df2.columns[[0]], axis=1, inplace=True)
```

```
[26]: df2.head(10)
```

	1	2	3	4	5	...	37	38	39	40	41
0	0	udp	private	SF	105	...	0.00.7	0.00.8	0.00.9	0.00.10	0.00.11
1	0	udp	private	SF	105	...	0.0	0.0	0.0	0.0	0.0
2	0	udp	private	SF	105	...	0.0	0.0	0.0	0.0	0.0
3	0	udp	private	SF	105	...	0.0	0.0	0.0	0.0	0.0
4	0	udp	private	SF	105	...	0.0	0.0	0.0	0.0	0.0
5	0	udp	private	SF	105	...	0.0	0.0	0.0	0.0	0.0
6	0	udp	domain_u	SF	29	...	0.0	0.0	0.0	0.0	0.0
7	0	udp	private	SF	105	...	0.0	0.0	0.0	0.0	0.0
8	0	udp	private	SF	105	...	0.0	0.0	0.0	0.0	0.0
9	0	tcp	http	SF	223	...	0.01	0.0	0.0	0.0	0.0

[10 rows x 41 columns]

```
[27]: ord_enc = OrdinalEncoder()
df2["make_tcp"] = ord_enc.fit_transform(df2[[2]])
df2[[2, "make_tcp"]].head(11)
```

	2	make_tcp
0	udp	2.0
1	udp	2.0
2	udp	2.0
3	udp	2.0
4	udp	2.0
5	udp	2.0
6	udp	2.0
7	udp	2.0
8	udp	2.0
9	tcp	1.0
10	udp	2.0

```
[28]: ord_enc = OrdinalEncoder()
df2["make_http"] = ord_enc.fit_transform(df2[[3]])
df2[[3, "make_http"]].head(11)
```



```
[28]:          3  make_http
0   private    46.0
1   private    46.0
2   private    46.0
3   private    46.0
4   private    46.0
5   private    46.0
6  domain_u    11.0
7   private    46.0
8   private    46.0
9     http     22.0
10  private    46.0
```

```
[29]: ord_enc = OrdinalEncoder()
df2["make_SF"] = ord_enc.fit_transform(df2[[4]])
df2[[4, "make_SF"]].head(11)
```

```
[29]:          4  make_SF
0   SF        9.0
1   SF        9.0
2   SF        9.0
3   SF        9.0
4   SF        9.0
5   SF        9.0
6   SF        9.0
7   SF        9.0
8   SF        9.0
9   SF        9.0
10  SF        9.0
```

```
[30]: df2.drop([2, 3, 4], axis=1, inplace=True)
df2.head(10)
```

```
[30]:    1     5     6     7     8  ...    40    41  make_tcp  make_http  make_SF
0  0  105  146  0.1  0.2  ...  0.00.10  0.00.11      2.0      46.0      9.0
1  0  105  146  0.0  0.0  ...      0.0      0.0      2.0      46.0      9.0
2  0  105  146  0.0  0.0  ...      0.0      0.0      2.0      46.0      9.0
3  0  105  146  0.0  0.0  ...      0.0      0.0      2.0      46.0      9.0
4  0  105  146  0.0  0.0  ...      0.0      0.0      2.0      46.0      9.0
5  0  105  146  0.0  0.0  ...      0.0      0.0      2.0      46.0      9.0
6  0   29     0  0.0  0.0  ...      0.0      0.0      2.0      11.0      9.0
7  0  105  146  0.0  0.0  ...      0.0      0.0      2.0      46.0      9.0
8  0  105  146  0.0  0.0  ...      0.0      0.0      2.0      46.0      9.0
9  0  223  185  0.0  0.0  ...      0.0      0.0      1.0      22.0      9.0
```

[10 rows x 41 columns]

```
[31]: df2 = df2.iloc[0:]
l=[]
for i in df2.iloc[0]:
```

```

    if len(str(i))>3:
        l.append(str(i)[0:4])
    else:
        l.append(i)
df2.iloc[0]=np.array(l)

```

```
[32]: df2.head(10)
```

```

[32]:   1    5    6    7    8    9  ...   39   40   41  make_tcp  make_http  make_SF
0  0  105  146  0.1  0.2  0.3  ...  0.00  0.00  0.00         2.0         46.0         9.0
1  0  105  146    0    0    0  ...  0.0   0.0   0.0          2          46          9
2  0  105  146    0    0    0  ...  0.0   0.0   0.0          2          46          9
3  0  105  146    0    0    0  ...  0.0   0.0   0.0          2          46          9
4  0  105  146    0    0    0  ...  0.0   0.0   0.0          2          46          9
5  0  105  146    0    0    0  ...  0.0   0.0   0.0          2          46          9
6  0   29    0    0    0    0  ...  0.0   0.0   0.0          2          11          9
7  0  105  146    0    0    0  ...  0.0   0.0   0.0          2          46          9
8  0  105  146    0    0    0  ...  0.0   0.0   0.0          2          46          9
9  0  223  185    0    0    0  ...  0.0   0.0   0.0          1          22          9

```

[10 rows x 41 columns]

```

[33]: X_test=df2.values
      print(len(X_test))

```

311029

```

[:]: #pca = PCA(n_components=10)
      #X_test=pca.fit_transform(X_test)

```

## 0.2 Applying Logistic regression Model

```

[:]: clf = LogisticRegression(penalty='l2',random_state=0,multi_class='auto').
      ↪fit(X_train, y_train)

```

```
[:]: y_predval = clf.predict(X_val)
```

```
[:]: precision_score(y_val, y_predval, average='weighted')
```

```
[:]: 0.9616212081493103
```

```
[:]: recall_score(y_val, y_predval, average='weighted')
```

```
[:]: 0.9678134368338157
```

```
[:]: f1_score(y_val, y_predval, average='weighted')
```

```
[:]: 0.9631995133442403
```

```
[:]: from sklearn.metrics import accuracy_score
```

```
accuracy_score(y_val, y_predval)
```

```
[ ]: 0.9678134368338157
[ ]: y_pred = clf.predict(X_test)
[ ]: prediction = pd.DataFrame(y_pred).to_csv('testLabelLog.csv',header=["target"])
[ ]: from google.colab import drive
[ ]: drive.mount('/content/drive')
[ ]: y_pred
[ ]: array(['normal.', 'normal.', 'normal.', ..., 'normal.', 'normal.',
          'normal.'], dtype=object)
```

### 0.3 Applying the SVM Model

```
[ ]: ylin_clf = svm.LinearSVC()
[ ]: lin_clf.fit(X_train, y_train)
[ ]: LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
             intercept_scaling=1, loss='squared_hinge', max_iter=1000,
             multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
             verbose=0)
[ ]: y_predval = lin_clf.predict(X_val)
[ ]: precision_score(y_val, y_predval, average='weighted')
[ ]: 0.9974327233474799
[ ]: recall_score(y_val, y_predval, average='weighted')
[ ]: 0.9939386890791236
[ ]: f1_score(y_val, y_predval, average='weighted')
[ ]: 0.9956167275161313
[ ]: from sklearn.metrics import accuracy_score
[ ]: accuracy_score(y_val, y_predval)
[ ]: 0.9939386890791236
[ ]: y_pred = lin_clf.predict(X_test)
[ ]: prediction = pd.DataFrame(y_pred).to_csv('testLabelSVM.csv',header=['target'])
[ ]: y_pred
[ ]: array(['normal.', 'normal.', 'normal.', ..., 'normal.', 'normal.',
          'normal.'], dtype=object)
[ ]: # Play an audio beep. Any audio URL will do.
[ ]: from google.colab import output
```

```
output.eval_js('new Audio("https://upload.wikimedia.org/wikipedia/commons/0/05/↪Beep-09.ogg").play()')
```

## 0.4 Applying NaiveBayes

```
[ ]: from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB().fit(X_train, y_train)

[ ]: y_predval = gnb.predict(X_val)

[ ]: precision_score(y_val, y_predval, average='weighted')

[ ]: 0.9904090357675394

[ ]: recall_score(y_val, y_predval, average='weighted')

[ ]: 0.9484435954127544

[ ]: f1_score(y_val, y_predval, average='weighted')

[ ]: 0.9658755438779745

[ ]: from sklearn.metrics import accuracy_score

accuracy_score(y_val, y_predval)

[ ]: 0.9484435954127544

[ ]: y_pred = gnb.predict(X_test)

[ ]: prediction = pd.DataFrame(y_pred).to_csv('testLabelGNB.csv',header=["target"])

[ ]: y_pred

[ ]: array(['normal.', 'normal.', 'normal.', ..., 'normal.', 'normal.',
        'normal.'], dtype='<U16')

[ ]: # Play an audio beep. Any audio URL will do.
from google.colab import output
output.eval_js('new Audio("https://upload.wikimedia.org/wikipedia/commons/0/05/↪Beep-09.ogg").play()')
```

## 0.5 Applying Random forests:

```
[34]: from sklearn.ensemble import RandomForestClassifier
print(len(X))
```

4898430

```
[35]: clf = RandomForestClassifier(n_estimators=500,max_depth=15, random_state=0)
clf.fit(X_train,y_train)
```

```
[35]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=15, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=500,
                             n_jobs=None, oob_score=False, random_state=0, verbose=0,
                             warm_start=False)

[37]: y_predval = clf.predict(X_val)

[38]: precision_score(y_val, y_predval, average='weighted')

[38]: 0.9999359544600412

[39]: recall_score(y_val, y_predval, average='weighted')

[39]: 0.9999387558908791

[40]: f1_score(y_val, y_predval, average='weighted')

[40]: 0.999936184951831

[41]: from sklearn.metrics import accuracy_score

      accuracy_score(y_val, y_predval)

[41]: 0.9999387558908791
```

### 0.5.1 Random forest produced the best results so far. So training with it to produce final test-Label.csv

```
[ ]: clf = RandomForestClassifier(n_estimators=500, max_depth=15, random_state=0)
      clf.fit(X, Y)

[ ]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=15, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=500,
                             n_jobs=None, oob_score=False, random_state=0, verbose=0,
                             warm_start=False)

[43]: y_pred = clf.predict(X_test)
      print(len(y_pred))
      print(len(X_test))
```

311029  
311029

```
[45]: prediction = pd.DataFrame(y_pred).
      ↳to_csv('submission_12_2020201012_2020201087_2020201095_2020202020.
      ↳csv',header=['target'])

[44]: y_pred

[44]: array(['normal.', 'normal.', 'normal.', ..., 'normal.', 'normal.',
          'normal.'], dtype=object)

[42]: # Play an audio beep. Any audio URL will do.
      from google.colab import output
      output.eval_js('new Audio("https://upload.wikimedia.org/wikipedia/commons/0/05/
      ↳Beep-09.ogg").play()')
```

## 1 Applying Decision Trees

```
[ ]: from sklearn.tree import DecisionTreeClassifier
      clf = DecisionTreeClassifier(random_state=0)
      clf.fit(X_train,y_train)

[ ]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                          max_depth=None, max_features=None, max_leaf_nodes=None,
                          min_impurity_decrease=0.0, min_impurity_split=None,
                          min_samples_leaf=1, min_samples_split=2,
                          min_weight_fraction_leaf=0.0, presort='deprecated',
                          random_state=0, splitter='best')

[ ]: y_predval = clf.predict(X_val)

[ ]: precision_score(y_val, y_predval, average='weighted')

[ ]: 0.9999953298485896

[ ]: recall_score(y_val, y_predval, average='weighted')

[ ]: 0.9999950509810811

[ ]: f1_score(y_val, y_predval, average='weighted')

[ ]: 0.9999951262046003

[ ]: from sklearn.metrics import accuracy_score

      accuracy_score(y_val, y_predval)

[ ]: 0.9999950509810811

[ ]: y_pred = clf.predict(X_test)

[ ]: prediction = pd.DataFrame(y_pred).
      ↳to_csv('submission_12_2020201012_2020201087_2020201095_2020202020_1.
      ↳csv',header=['target'])
```

```
[ ]: y_pred
```

```
[ ]: array(['normal.', 'normal.', 'normal.', ..., 'normal.', 'normal.',  
         'normal.'], dtype=object)
```

```
[47]: %%capture
```

```
!wget -nc https://raw.githubusercontent.com/brpy/colab-pdf/master/colab_pdf.py  
from colab_pdf import colab_pdf  
colab_pdf('SNS_ASS5.ipynb')
```

```
[ ]:
```